


Using AADL for Microarchitecture Models: Initial Experiences & Suggestions


Mark Brown
Lynx Software Technologies



Make it simple.
Make it last.




US Navy




Boeing

Industries Served



Wabtec




Elekta

Security


- Military
- Government
- Critical National Infrastructure
- Enterprise
- FinTech
- IIoT

Safety


- Avionics
- Automotive
- Railway
- Industrial Control
- Medical



Lockheed Martin



Boeing

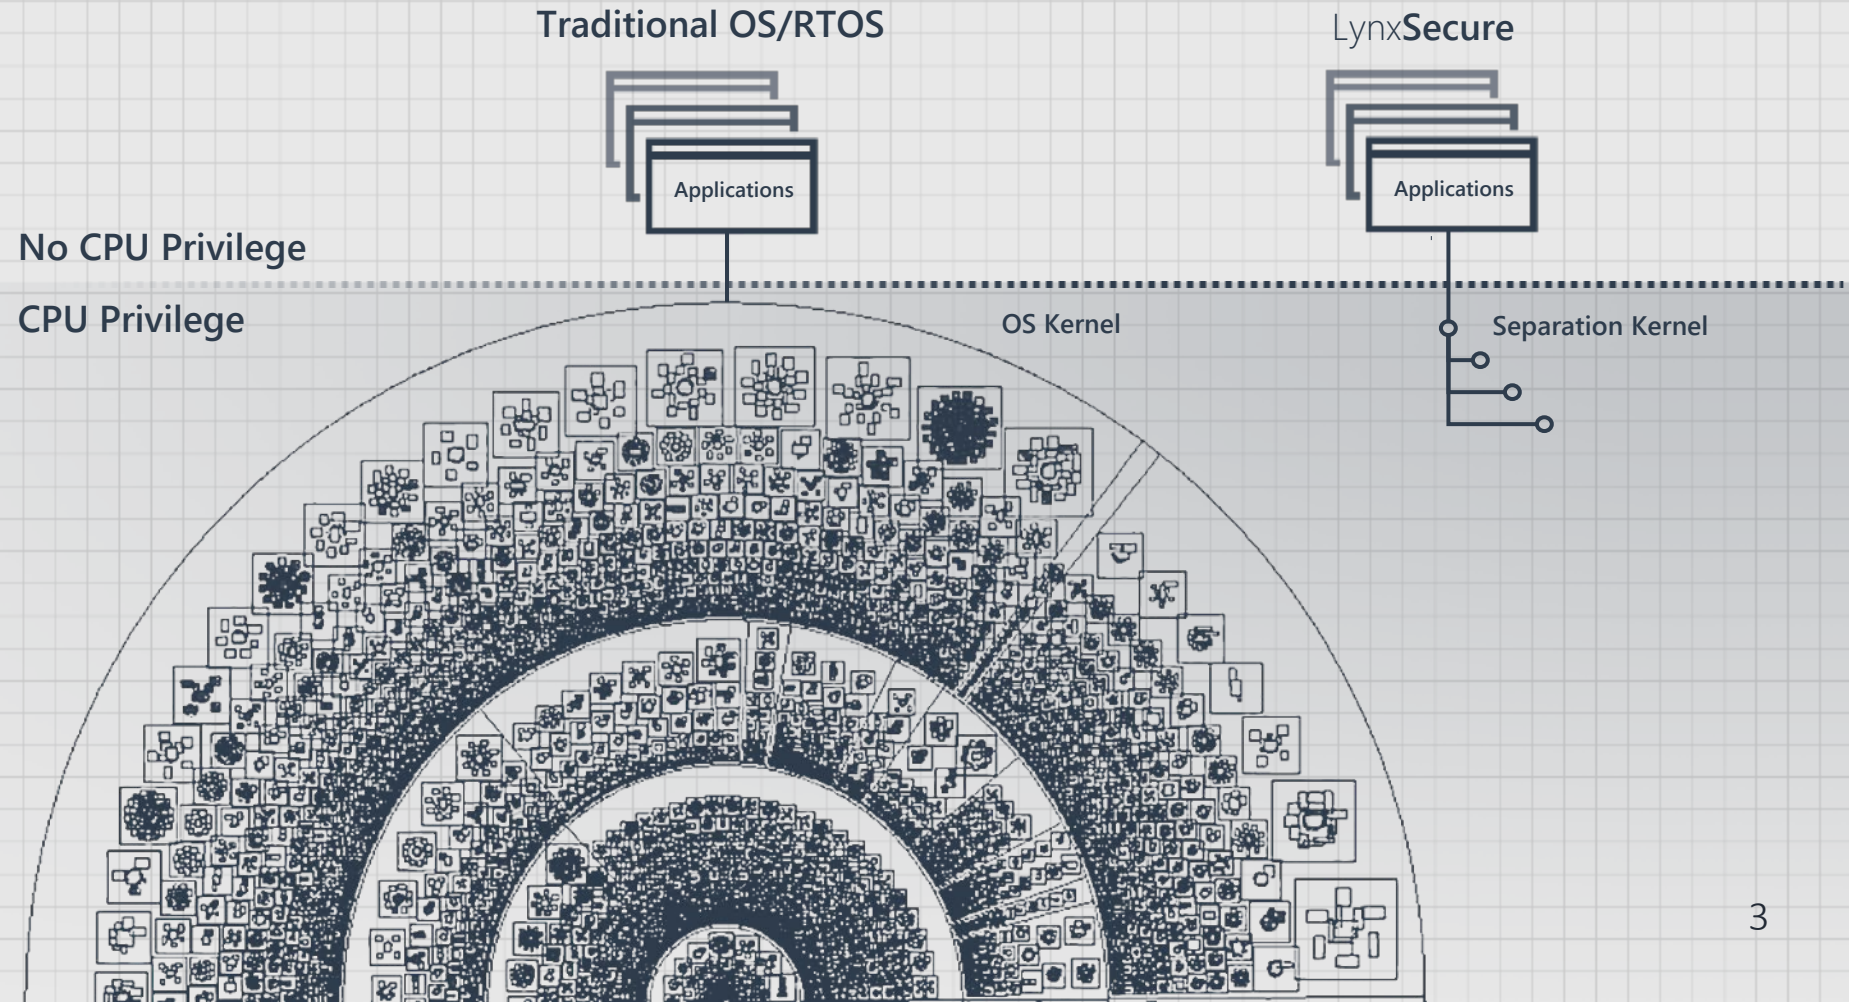


Bosch + ETAS

Assurance



"Kernel" Comparison









The answer lies
within
Architecture


Software Architecture has the greatest influence on:

- Safety and Security Properties
- Performance
- System Comprehensibility
- Design Adaptability & Reuse
- Development Efficiency
- Novice Adoption



Product Technologies

Hardware Virtualization




LYNX | Secure®

RTOS



LYNX | OS-178®

Modular Development Framework

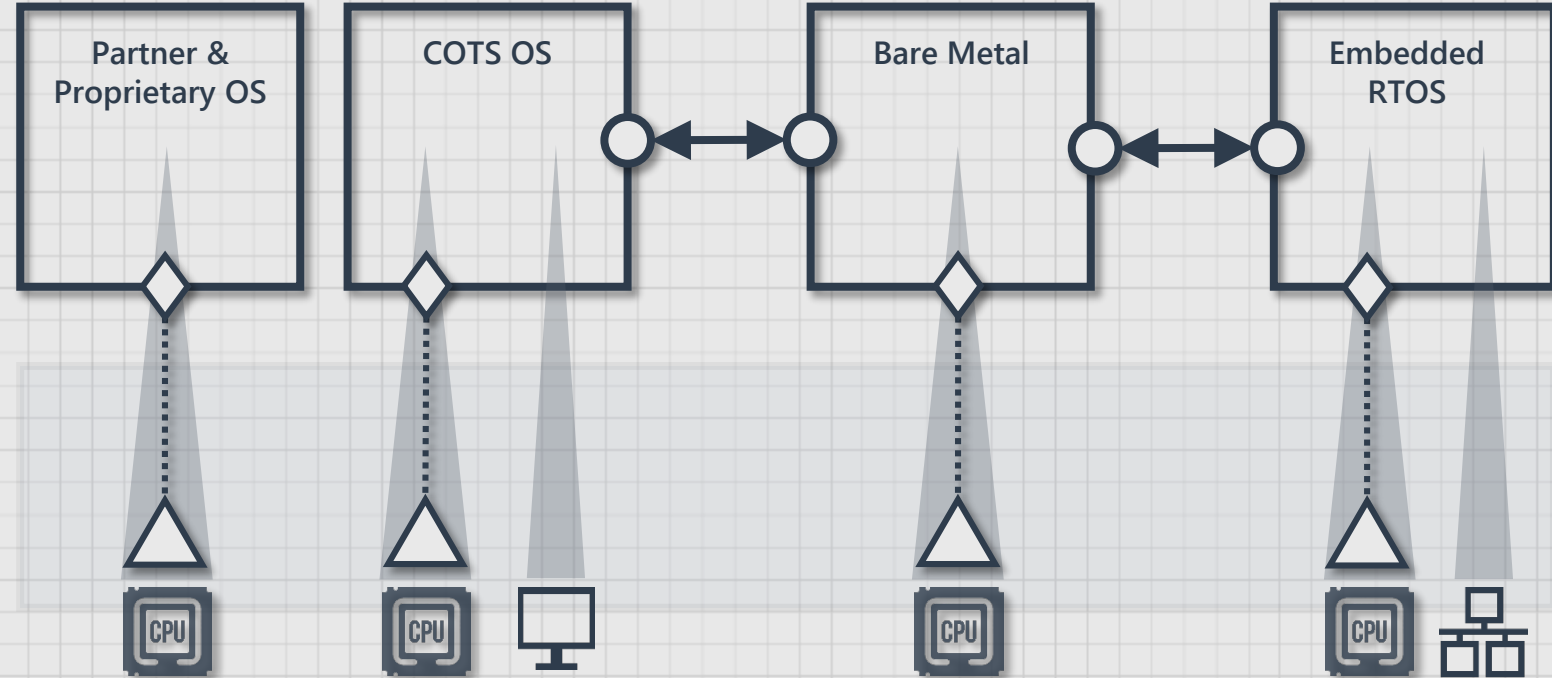


LYNX
MOSA.ic™

LYNX Simplified Modular Integration Framework



- **Rooms** (AADL 'system' – with 'processor' / 'virtual processor', IO 'device', 'process' &c.)
- **Core Software** (An OS can be "virtualized" to run as a Guest Operating System)
- **Passageways** (lock-free communications for 'port', e.g., ARINC653 protocols)
- **IO/Resource Assignment** (static 'binding' of hardware resources to each Room/'system')



AADL model of *uHardware*?

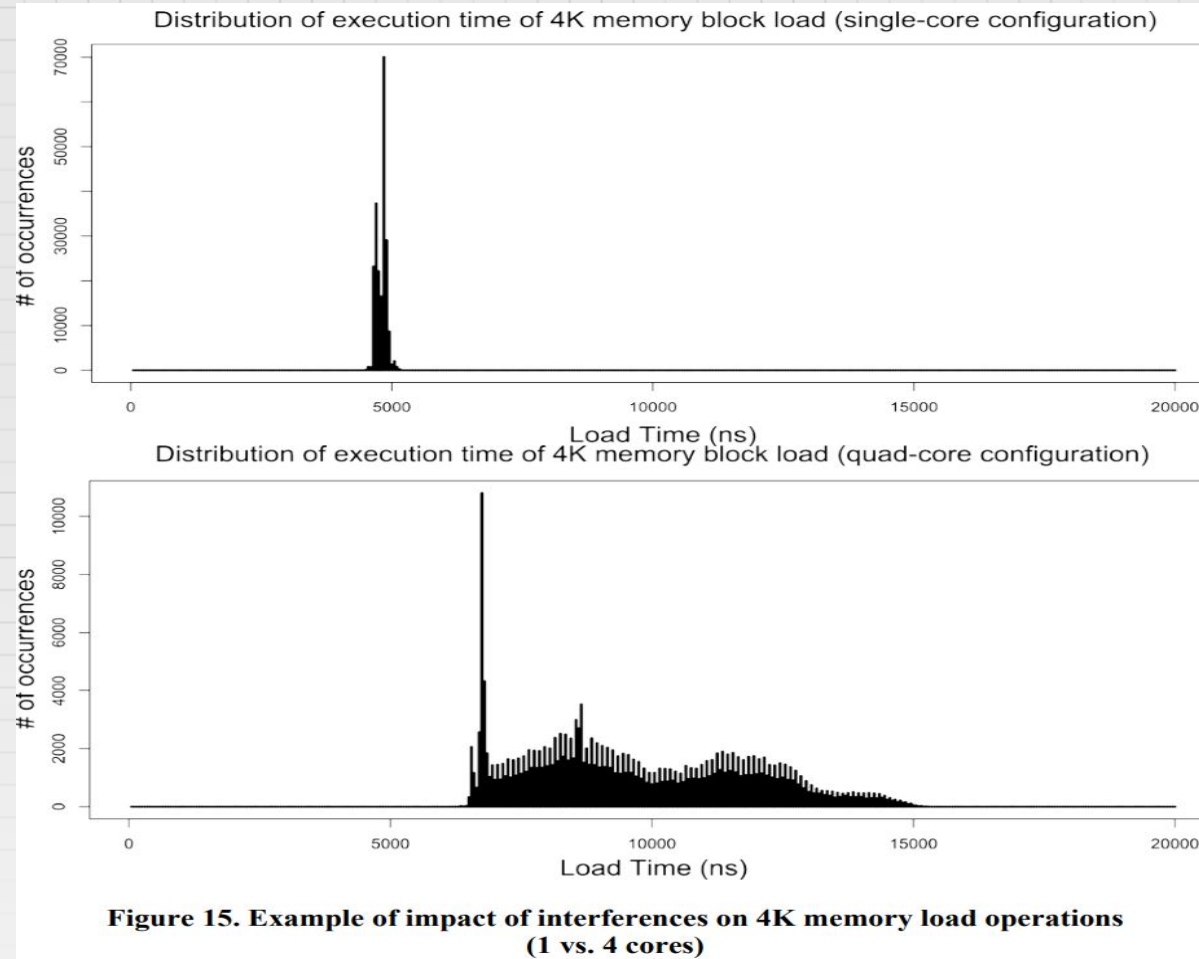
- Trying out AADL/ReqSpec for certification
- Can it express detail for MMU hardware?
 - Isn't that just a 'device'?
 - With a 'bus' and a 'processor'?
 - Can I build a 'port' for pagetable 'data'?
 - Etc.
- Asked two more questions
 - MultiCore Processor (CAST-32A):
"Hardware Interference Channels"?
 - Fetching and Speculative Execution (Spectre):
"Hardware bugs?"
- Want a model for *micro-architecture*



MCP Problem Focus

WCET *despite* Memory Latency Non-Determinism

SCP: Tightly
bunched histogram




MCP: Widely spread
histogram
(non-deterministic)


*“For both cases, we collect memory access time and application’s execution time **while one core executes the benchmark and others execute a stressing benchmark over the same DRAM controller.**”*
(FAA TC-16/51, p. 48)

What is the MCP Problem?

A Better View



Single Core Processor



Multi-Core Processor

- Complex interconnect
 - Roundtable
 - Protocols, queues, caches, expensive, growing
 - Congestion
 - Unpredictable initiations
 - Uncontrolled initiations
- Memory competition
 - Active cores
 - Active DMA
 - Preferred lanes

MCP uHW Proposed Model (Evenblij 2017)


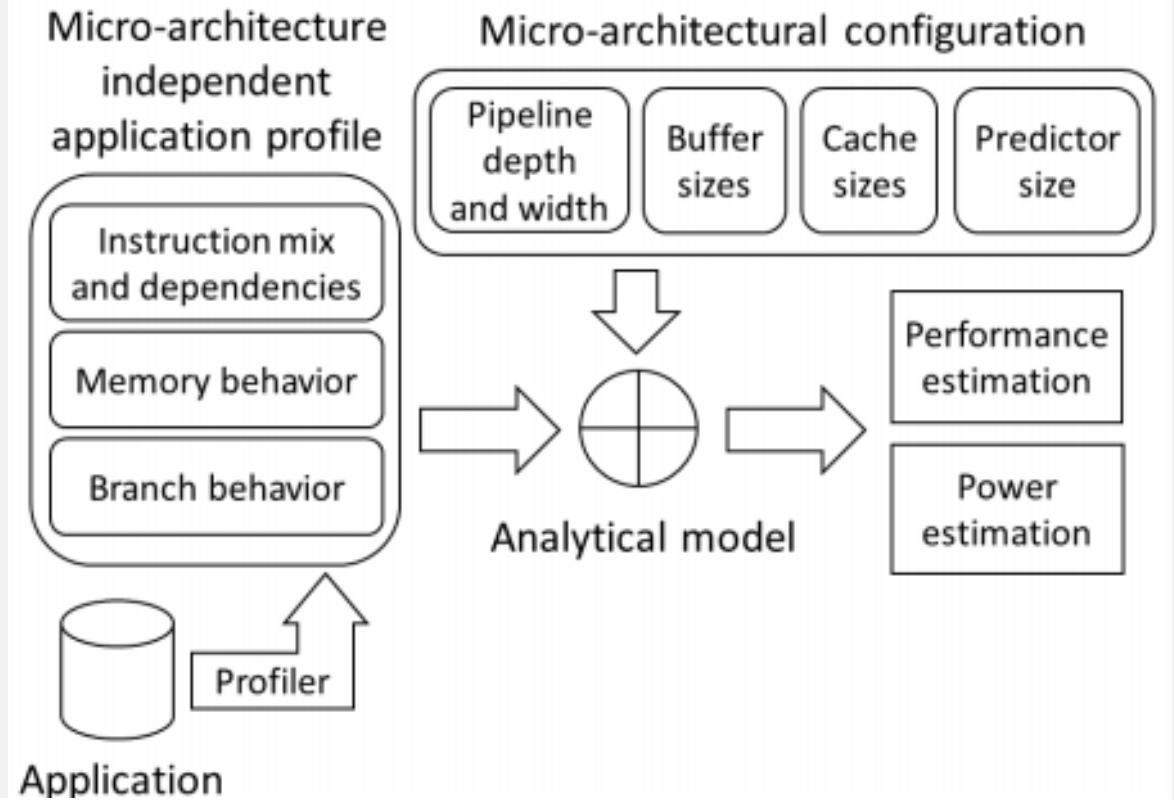
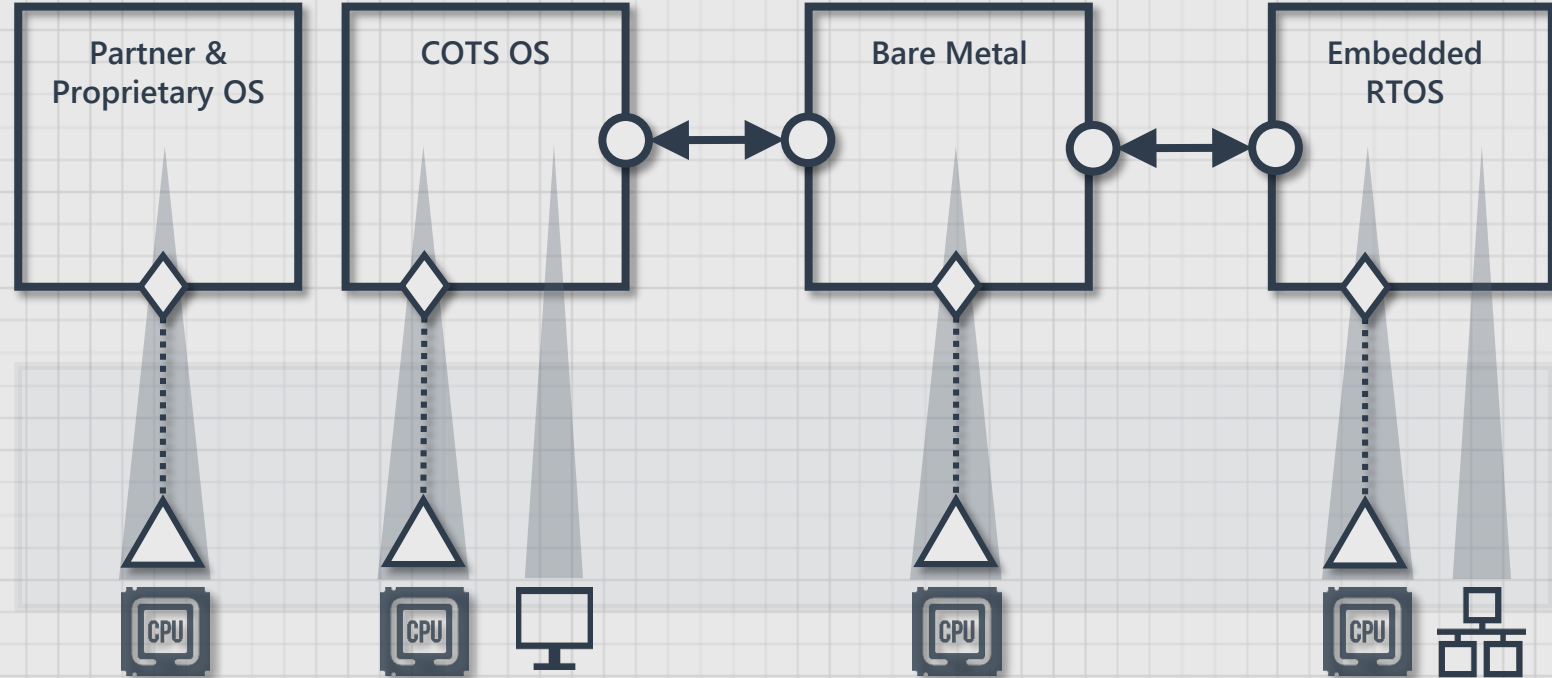


Figure 2.2: Interval analysis divides execution time in intervals between miss events [1].

LYNX Simplified Modular Integration Framework



- **Rooms** (AADL 'system' – with 'processor' / 'virtual processor', IO 'device', 'process' &c.)
- **Core Software** (An OS can be "virtualized" to run as a Guest Operating System)
- **Passageways** (lock-free communications for 'port', e.g., ARINC653 protocols)
- **IO/Resource Assignment** (static 'binding' of hardware resources to each Room/'system')

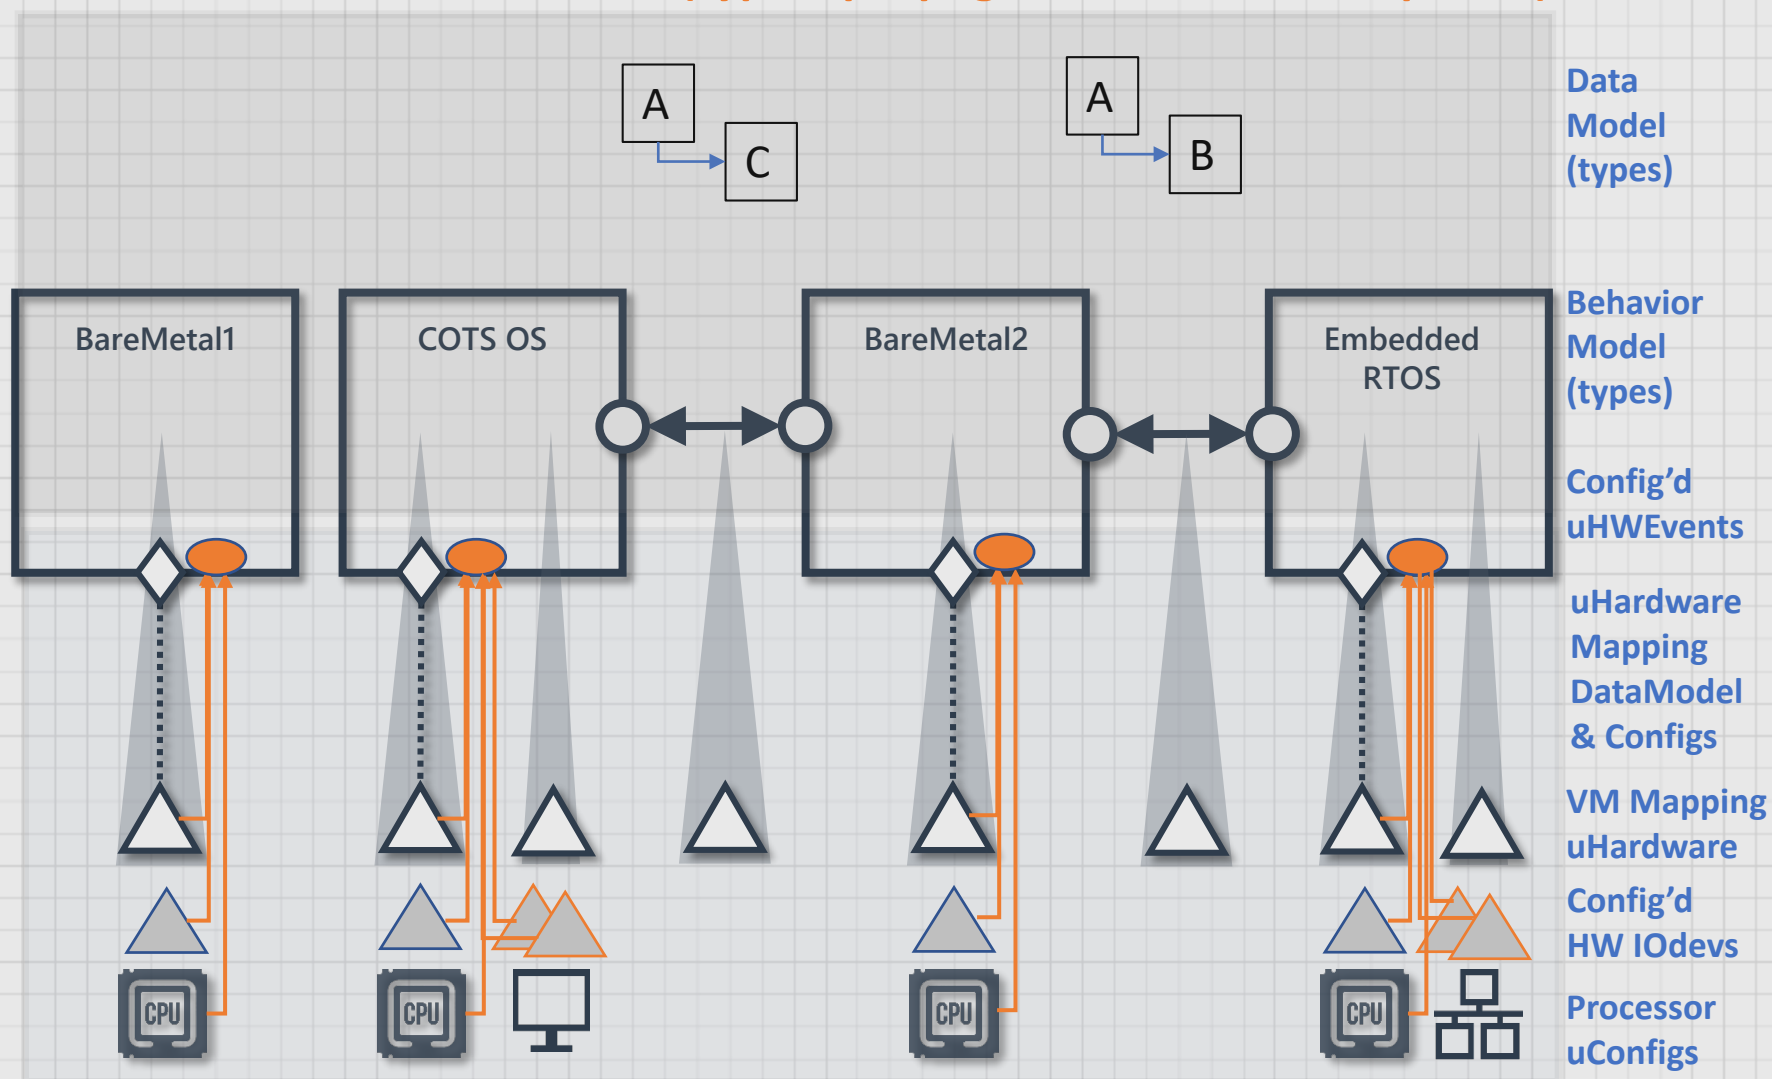




LYNX Simplified Modular Integration **uHardware** Framework



Health Monitor / Error Model (types, propagation, behavior, repair,...)



AADL model of *uHardware*?

- Two questions
 - MultiCore Processor (CAST-32A):
"Hardware Interference Channels"?
 - Spectre:
"Hardware bugs?"
- AADL *micro-architecture* model:
 - 'system' – one per Room; configured container
 - 'processor' – "contains" nearly all of the problems
 - 'bus' – use these to compose uHW into graph (topology)
 - 'device' – uHW; composed using buses
 - 'data' – used to configure uHW, e.g., pagetable
- *Not* HDL code (right?)

Initial Requirements (1)

- CAST-32A Hardware Inteference
 - Model a *HW Topology* for memory access
 - from processor core
 - through caches
 - through interconnect (cache coherent fabric)
 - through chokepoints
 - to DRAM device
 - Model the *memory access latency*
 - Cache “hit” means good; negative latency
 - Cache “miss” bad; latency added
 - Latency accumulates per node walked in the graph!
 - Model latency as EMv2 '*error*'
 - errors and data 'propagate' via a 'port'
 - total accumulation of latency (store/load round-trip)
 - from request, through HW Topology, to satisfy (port)

Initial Requirements (2)

- Spectre Hardware Bug (Kocher et al.)
 - Model a *HW Topology* for memory access
 - in a processor core, chip-specific
 - via core-local buses
 - including all relevant core-local uHW
 - query for sufficient conditions-to-trigger
 - Model the *uHW that fails to perform access check*
 - We've landed in "reverse engineering" territory!
 - (i.e., Google, Kocher, et al. found a fault / root cause)
 - Increasingly need to collaborate with origin labs
 - Increasingly need to collaborate with chip makers
 - Model uHW conditions that activate / trigger the fault
 - Some in uHW, e.g., TLB / Cache, Speculative Execution Unit..
 - Some in system-level software state, e.g., MMU page tables



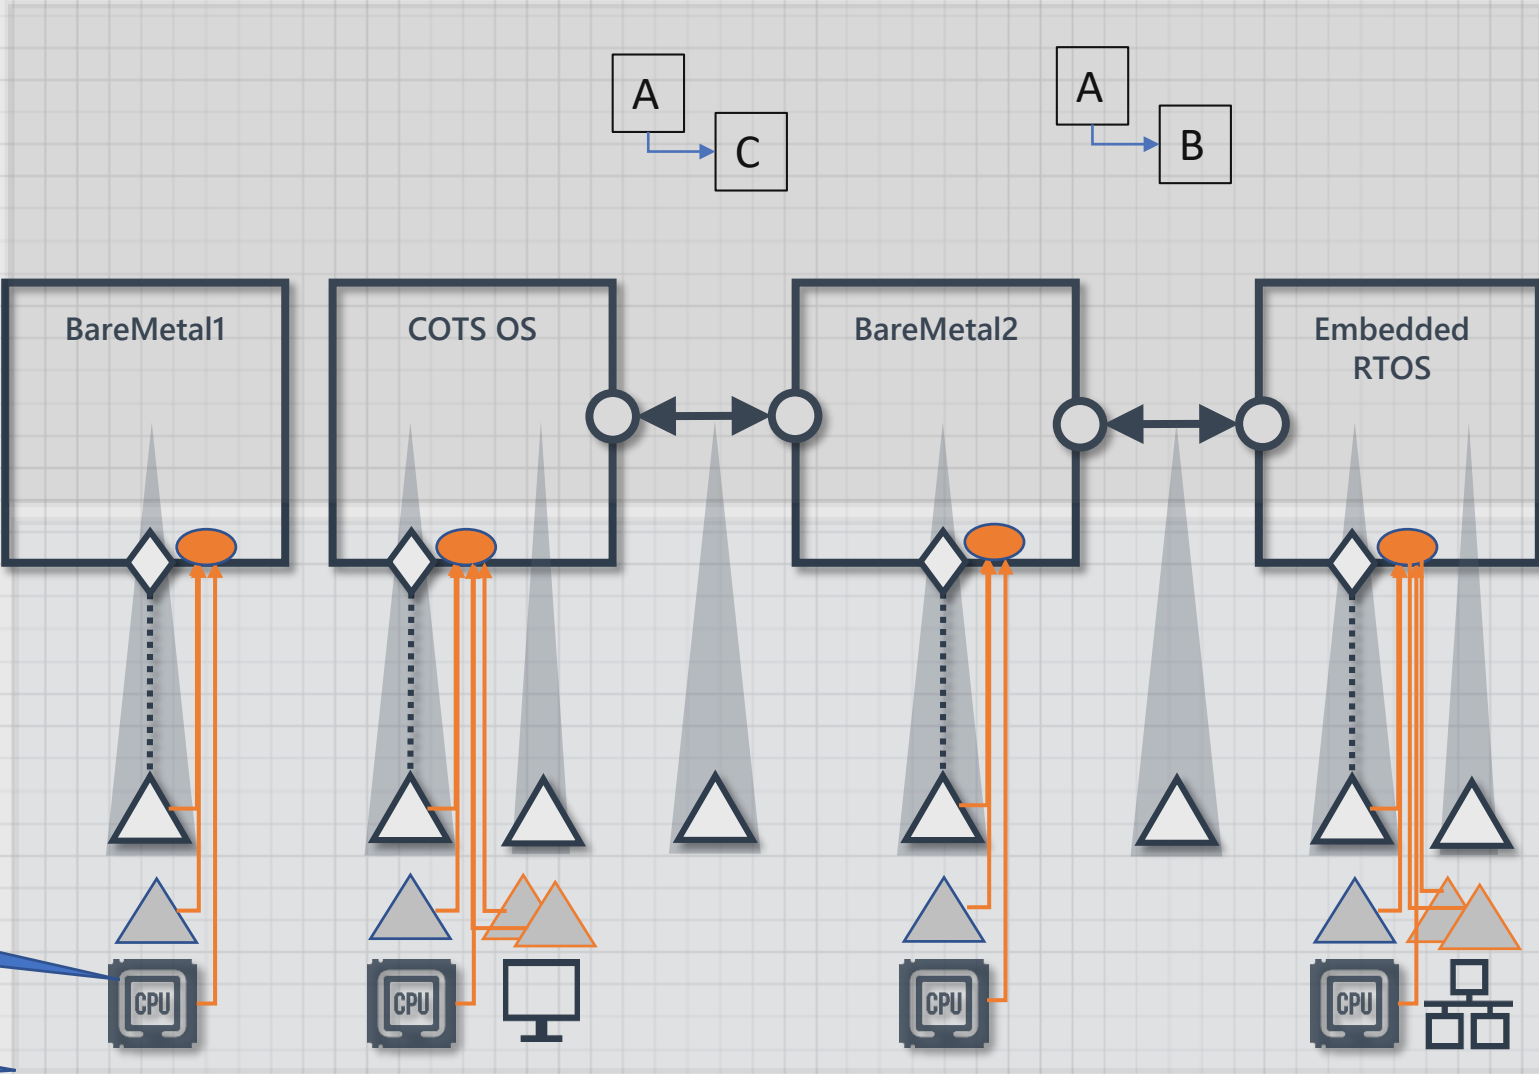
Health Monitor / Error Model (types, propagation, behavior, repair,...)

LYNX's
uHardware
AADL Library



Spectre &c.

CAST-32A



Data
Model
(types)

Behavior
Model
(types)

Config'd
uHWEvents

uHardware
Mapping
DataModel
& Configs

VM Mapping
uHardware

Config'd
HW IOdevs

Processor
uConfigs

Observations (1)

What's the *purpose* for uHW modeling?

- Is AADL for Teamwork or for Analysis?
 - published model (from chipmaker TRM)
 - private, reverse-engineered model (probably not right)
 - private AADL model *query*?
 - Observation: AADL is a standard, which promotes **teams**
- How can AADL standards promote teamwork?
- Does the AADL System Integrator want:
 - Better hardware models?
 - Hardware-resource error models?
 - Pre-built certification templates?
 - Can these expensive "options" align with "faster & cheaper?"

Observations (2)

What's the "Customer Demand" (for a simple, modular, AADL uHW Library)?

- "Do it for me" – wants to buy the uHW model
 - uHW AADL 'device', 'bus', 'data'/Data model, 'ports', etc.
 - "Just to look at" – to analyze, query, & re-use
 - "I don't want to pay a big AADL cost-to-entry"
- "Hold the errors" – doesn't want any!
 - Some errors are EMv2-repaired ('error sink')
 - Some errors always 'propagate'
 - Some 'error behavior' should be extensible/overridable
 - Observation: hardware resources as a primary 'error source' (e.g., 'ARINC653::HM_Error_ID_Actions', 'Error_Level_Type', etc.)
- Transparency – about our LynxSecure requirements
 - SW: validation needs to know how well it works (i.e., the partitioning architecture's internal requirements)
 - HW: hardware needs software to register the event handler (as an integration requirement)
 - Data: A customer can configure partitioning for safety – *or not!* (i.e., configuration data could specify poor requirements)

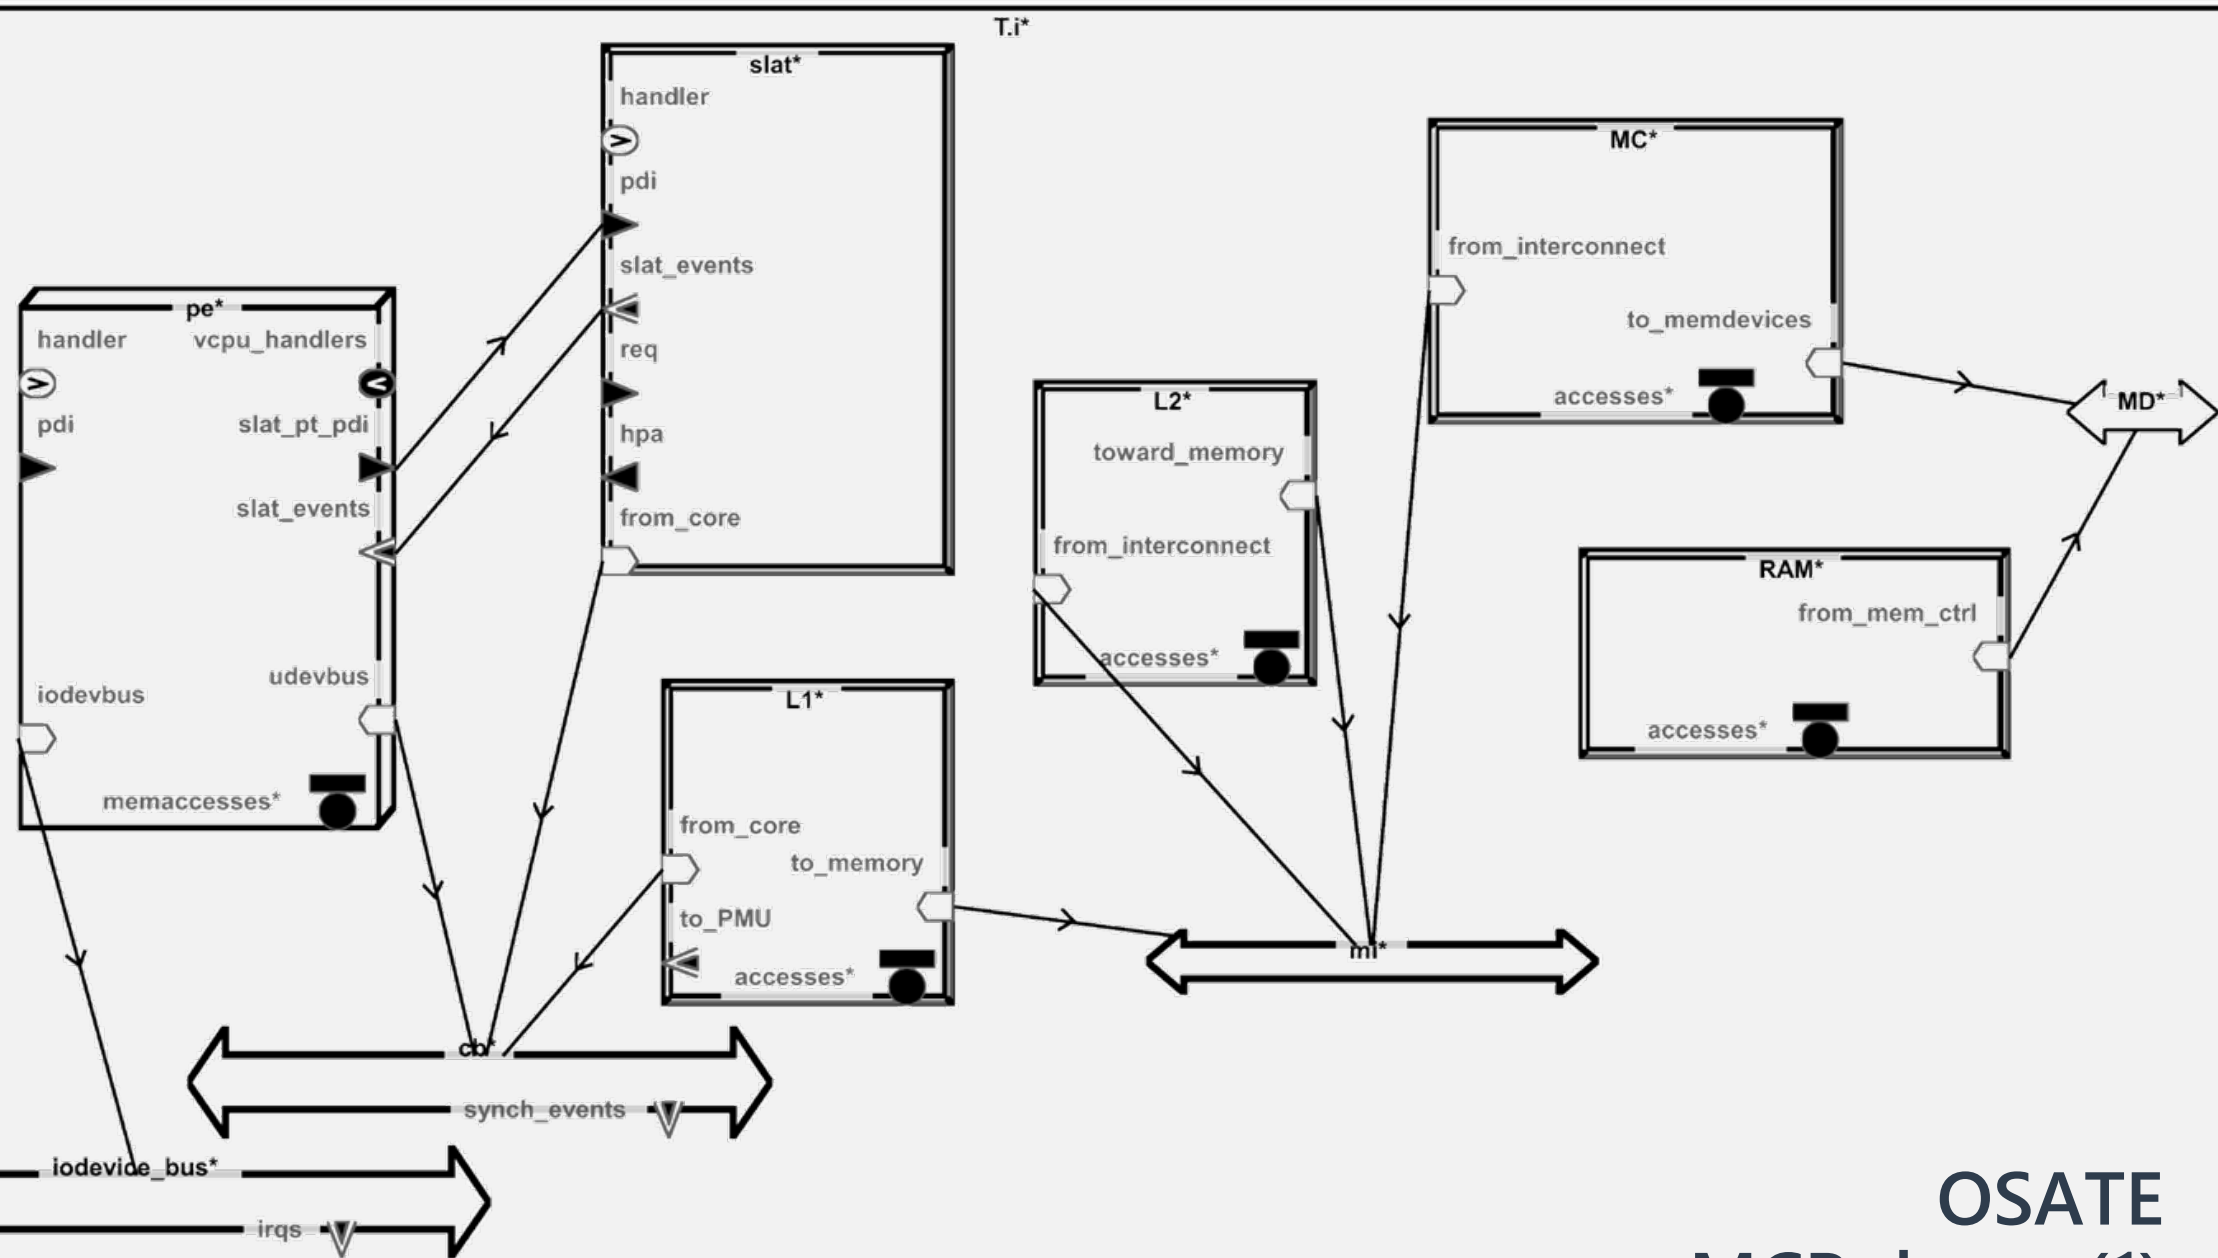
Opportunities/ Current Goals

Best practices

- OSTATE provides formal methods
- Requirements engineering?
- Certification templates?

Prebuilt AADL models “below the HAL” (uHW)

- Prebuilt repository for AADL-style requirements & model
- Prebuilt models for uHW resources
- Model for MCP partitions
- Working error model



OSATE MCP demo (1)

OSATE demo (2) "partitioning/"

- Abstract Model
 - Hardware (hw)
 - Software (handlers)
 - Data (pdi)
 - Hardware abstraction layer (resources)
 - Abstract HW concepts (control_type)
- ReqSpec requirements
 - for Hardware
 - for Software
 - for Data
 - to define/explain concepts
 - control_type
 - resources

OSATE demo (3) "aarch64/"

- HW-specific model (aarch64)
 - hw
 - handlers
 - pdi
 - resources
- ReqSpec for aarch64
 - As before: {hw, handlers, pdi}
 - This time, traces to ARM TRMs
 - This time, defines "resources" precisely, in terms of HW TRM definitions
 - Defines what a configured 'system' means

OSATE demo

What's missing? (1)

Automated AADL analysis for MCP?

- **Start:** data used to configure uHW blocks
- **Add:** hardware reverse engineering (chip vendor)
 - This is expensive; I cannot share
 - Basic idea is to measure each uHW block
 - What can possibly trigger latency?
 - What's the latency cost, per uHW block?
 - What's the worst case (consecutive losses...)
 - Goal is to measure all uHW blocks (AADL properties)
- **Add:** software memory use metrics (customer)
 - Record memory use profile per thread/SW
 - Find the worst possible competitions, given uHW + SW
- **Deploy:** HW interference mitigations (HW+SW+data)
 - Break up the fights! (dynamic, not AADL-static)
 - Prove that the new worst case is better (how much?)

What's missing? (2)

AADL analysis for MCP – by parts

- **Start:** data used to configure uHW blocks
- **Added:** hardware reverse engineering (chip vendor)
- **Measured:** software memory use metrics (customer)
 - Record memory use profile per thread/SW
 - Find the worst possible competitions, given uHW + SW
- **Modeled:** HW interference mitigations (HW+SW+data)
 - Break up the fights! (dynamic, not AADL-static)
 - Estimate: mitigated case is better (how much?)

Suggestions for AADL Committee (1)

Anticipate Pre-built AADL models

- concept: AADL libraries, with reqspec, etc.
- customer “just wants to build/integrate software”
- reusable by AADLv3 Configuration
- e.g., ‘Customer1 system extends Lynx::Room’
 - was ‘Customer1 system’ (no extends)
 - or was ‘Customer1 system extends Other::System’
- e.g., ‘message1 data extends balsa_data_model::EGI_Data_Platform’

Suggestions for AADL Committee (2)

Anticipate AADL models for HW uArchitecture

- concept: AADL libraries, with reqspec, etc.
- customer “just wants to build/integrate software”
- library provides a model for some questions
 - reusable by inheritance
 - query-able (new formal methods queries)
 - configure-able (AADL configuration to quickly represent a chip or a target board or larger system, but with HiFi uHW)
- Promote teaming
 - reverse engineering of realistic chip models
 - collaboration with HW chip manufacturers, “Please confirm..”
 - encode best practices for/toward solutions
- Collaboration via git – or just a zipfile/tarball

Suggestions for AADL Committee (3)

Anticipate AADL models as certification repositories

- concept: AADL libraries, with reqspec, etc.
- customer wants the least possible:
 - certification headaches (examples, Configuration, reports)
 - HW / uHW / MCP headaches
 - customer just wants to build software
- library provides a template-model for best practices
 - reusable by structure (directories, packages, etc.)
 - per regime / per set of best practices (Agile-style?)
 - Architecture-led
- Promote education / learning
 - encode best practices for/toward solutions
 - some learn best from an example
- Collaboration via git – or just a zipfile/tarball

Suggestions for AADL Committee (4)

How best to model uHW?

- New 'device' – inside / outside of 'processor'?
 - All memory uHW affects every 'thread', 'process', etc.
 - Every uHW event switches the thread SW
 - A context switch to OS always pollutes the caches
 - A context switch from OS to not-me pollutes caches
- concept: AADL "zoom"
 - out: "just a processor"
 - in: uHW exploded view (breakdown into components)
- virtual vs. (not) processor / bus / memory

Call to Action

- Do you agree that:
 - AADL can/should be used to model *Hardware microarchitecture*?
 - The AADL standards, together with reqspec & friends, can be used to promote teamwork?
 - Teamwork will help with thorny issues like safety, security, reliability?
- How can you / your organization help?
 - Not everyone has expertise here
 - One solution may not fit all
 - Developing world-class exemplars may help most
- Do you want to
 - team with Lynx as we build up some exemplars?
 - obtain rich, detailed, automatable models of uHW?

Thank You.